

OUTPUTS:

Multiple regression model:

- Data preprocessing:

```
> head(dia_dummy)
  id carat cutGood cutIdeal cutPremium cutVery Good colorE colorF colorG colorH colorI colorJ clarityIF
1  1 0.23      0      1      0      0      0      1      0      0      0      0      0      0
2  2 0.21      0      0      1      0      0      1      0      0      0      0      0      0
3  3 0.23      1      0      0      0      0      1      0      0      0      0      0      0
4  4 0.29      0      0      1      0      0      0      0      0      1      0      0      0
5  5 0.31      1      0      0      0      0      0      0      0      0      1      1      0
6  6 0.24      0      0      0      0      1      0      0      0      0      0      1      0
  claritySI1 claritySI2 clarityVS1 clarityVS2 clarityVVS1 clarityVVS2 depth table price  x    y    z
1          0          1          0          0          0          0      61.5   55   326 3.95 3.98 2.43
2          1          0          0          0          0          0      59.8   61   326 3.89 3.84 2.31
3          0          0          1          0          0          0      56.9   65   327 4.05 4.07 2.31
4          0          0          0          1          0          0      62.4   58   334 4.20 4.23 2.63
5          0          1          0          0          0          0      63.3   58   335 4.34 4.35 2.75
6          0          0          0          0          0          1      62.8   57   336 3.94 3.96 2.48
```

- Feature selection:

Backward using p value elimination:

```
coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    2086.771     438.256   4.762 1.93e-06 ***
train.data$carat 11256.528     54.476 206.631 < 2e-16 ***
train.data$cutGood  588.726     37.650  15.637 < 2e-16 ***
train.data$cutIdeal 838.045     37.386  22.416 < 2e-16 ***
train.data$cutPremium 776.079     36.092  21.503 < 2e-16 ***
train.data$cutVery Good 736.480     36.098  20.402 < 2e-16 ***
train.data$colorE -212.232     20.076 -10.571 < 2e-16 ***
train.data$colorF -279.320     20.270 -13.780 < 2e-16 ***
train.data$colorG -499.123     19.882 -25.104 < 2e-16 ***
train.data$colorH -988.822     21.199 -46.644 < 2e-16 ***
train.data$colorI -1474.757     23.709 -62.203 < 2e-16 ***
train.data$colorJ -2380.119     29.229 -81.430 < 2e-16 ***
train.data$clarityIF 5383.561     57.196  94.125 < 2e-16 ***
train.data$claritySI1 3669.275     48.887  75.057 < 2e-16 ***
train.data$claritySI2 2697.368     49.087  54.950 < 2e-16 ***
train.data$clarityVS1 4578.481     49.902  91.749 < 2e-16 ***
train.data$clarityVS2 4268.354     49.133  86.873 < 2e-16 ***
train.data$clarityVVS1 5001.054     52.898  94.541 < 2e-16 ***
train.data$clarityVVS2 4960.726     51.372  96.564 < 2e-16 ***
train.data$depth   -61.094        4.583 -13.332 < 2e-16 ***
train.data$table    -27.738         3.274  -8.472 < 2e-16 ***
train.data$x      -1028.854       23.041 -44.653 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1134 on 43130 degrees of freedom
Multiple R-squared:  0.9192,    Adjusted R-squared:  0.9192
F-statistic: 2.338e+04 on 21 and 43130 DF,  p-value: < 2.2e-16
```

Backward using aic

```

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    2086.771    438.256   4.762 1.93e-06 ***
train.data$carat    11256.528    54.476 206.631 < 2e-16 ***
train.data$cutGood     588.726    37.650  15.637 < 2e-16 ***
train.data$cutIdeal     838.045    37.386  22.416 < 2e-16 ***
train.data$cutPremium    776.079    36.092  21.503 < 2e-16 ***
train.data$`cutVery Good`  736.480    36.098  20.402 < 2e-16 ***
train.data$colorE    -212.232    20.076 -10.571 < 2e-16 ***
train.data$colorF    -279.320    20.270 -13.780 < 2e-16 ***
train.data$colorG    -499.123    19.882 -25.104 < 2e-16 ***
train.data$colorH    -988.822    21.199 -46.644 < 2e-16 ***
train.data$colorI   -1474.757    23.709 -62.203 < 2e-16 ***
train.data$colorJ   -2380.119    29.229 -81.430 < 2e-16 ***
train.data$clarityIF    5383.561    57.196  94.125 < 2e-16 ***
train.data$claritySI1    3669.275    48.887  75.057 < 2e-16 ***
train.data$claritySI2    2697.368    49.087  54.950 < 2e-16 ***
train.data$clarityVS1    4578.481    49.902  91.749 < 2e-16 ***
train.data$clarityVS2    4268.354    49.133  86.873 < 2e-16 ***
train.data$clarityVVS1    5001.054    52.898  94.541 < 2e-16 ***
train.data$clarityVVS2    4960.726    51.372  96.564 < 2e-16 ***
train.data$depth     -61.094     4.583 -13.332 < 2e-16 ***
train.data$table     -27.738     3.274  -8.472 < 2e-16 ***
train.data$x       -1028.854    23.041 -44.653 < 2e-16 ***
---
signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1134 on 43130 degrees of freedom
Multiple R-squared:  0.9192,    Adjusted R-squared:  0.9192
Forward using aic:

```

Coefficients:



	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2086.771	438.256	4.762	1.93e-06	***
train.data\$carat	11256.528	54.476	206.631	< 2e-16	***
train.data\$claritySI2	2697.368	49.087	54.950	< 2e-16	***
train.data\$colorJ	-2380.119	29.229	-81.430	< 2e-16	***
train.data\$claritySI1	3669.275	48.887	75.057	< 2e-16	***
train.data\$colorI	-1474.757	23.709	-62.203	< 2e-16	***
train.data\$x	-1028.854	23.041	-44.653	< 2e-16	***
train.data\$colorH	-988.822	21.199	-46.644	< 2e-16	***
train.data\$depth	-61.094	4.583	-13.332	< 2e-16	***
train.data\$table	-27.738	3.274	-8.472	< 2e-16	***
train.data\$clarityVVS2	4960.726	51.372	96.564	< 2e-16	***
train.data\$clarityIF	5383.561	57.196	94.125	< 2e-16	***
train.data\$clarityVVS1	5001.054	52.898	94.541	< 2e-16	***
train.data\$clarityVS1	4578.481	49.902	91.749	< 2e-16	***
train.data\$clarityVS2	4268.354	49.133	86.873	< 2e-16	***
train.data\$colorG	-499.123	19.882	-25.104	< 2e-16	***
train.data\$colorF	-279.320	20.270	-13.780	< 2e-16	***
train.data\$colorE	-212.232	20.076	-10.571	< 2e-16	***
train.data\$cutIdeal	838.045	37.386	22.416	< 2e-16	***
train.data\$cutPremium	776.079	36.092	21.503	< 2e-16	***
train.data\$`cutVery Good`	736.480	36.098	20.402	< 2e-16	***
train.data\$cutGood	588.726	37.650	15.637	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1134 on 43130 degrees of freedom

Multiple R-squared: 0.9192, Adjusted R-squared: 0.9192

- Residual analysis:



- Postprocessing
Vif:

```
> vif(dia_back1)
      train.data$carat      train.data$cutGood      train.data$cutPremium      train.data$`cutVery Good`
      1.247245              1.306589              1.556387              1.348212
      train.data$colorE      train.data$colorF      train.data$colorG      train.data$colorH
      1.497443              1.479249              1.555223              1.429105
      train.data$colorJ      train.data$clarityIF      train.data$claritySI2      train.data$clarityVS1
      1.201697              1.145108              1.417407              1.387752
      train.data$clarityVS2      train.data$clarityVVS1      train.data$clarityVVS2      train.data$depth
      1.480602              1.244905              1.296324              1.179342
      train.data$table
      1.549290
```

Influential points:

```
cooksdl <- cooks.distance(dia_back1)
# influential row numbers
influential1 <- as.numeric(names(cooksdl)[(cooksdl > 4/nrow(train.data))])
head(train.data[influential1, ])
train.data1 <- train.data[-c(19534, 31109, 17573, 37886, 32318), ]
```

	carat	cutGood	cutIdeal	cutPremium	cutVery Good	colorE	colorF	colorG	colorH
52533	0.70	0	1	0	0	0	0	0	0
8379	0.31	0	1	0	0	0	1	0	0
17651	1.08	0	1	0	0	0	1	0	0
52971	0.40	0	1	0	0	0	0	0	0
41013	0.50	0	1	0	0	1	0	0	0
38591	0.57	0	0	0	1	0	0	0	0
	colorI	colorJ	clarityIF	claritySI1	claritySI2	clarityVVS1	clarityVVS2	clarityVVS1	
52533	1	0	0	0	0	0	0	0	0
8379	0	0	0	0	0	0	0	1	0
17651	0	0	0	0	0	0	0	1	0
52971	0	1	0	1	0	0	0	0	0
41013	0	0	0	0	1	0	0	0	0
38591	0	1	0	1	0	0	0	0	0
	clarityVVS2	depth	table	price	x	y	z		
52533	1	61.1	56	2530	5.73	5.76	3.52		
8379	0	61.9	57	583	4.32	4.34	2.68		
17651	0	61.9	55	7110	6.58	6.64	4.09		
52971	0	62.2	57	552	4.71	4.74	2.94		
41013	0	61.1	58	1185	5.09	5.13	3.12		
38591	0	60.6	61	1037	5.36	5.40	3.26		

- Final model:

```

coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    5727.729    396.106   14.460 < 2e-16 ***
train.data1$carat    8615.753    14.579  590.956 < 2e-16 ***
train.data1$cutGood   -38.400    24.595   -1.561  0.118
train.data1$cutPremium    95.585    17.699    5.401 6.67e-08 ***
train.data1$`cutVery Good`  90.961    17.235    5.278 1.31e-07 ***
train.data1$colorE    376.604    19.619   19.196 < 2e-16 ***
train.data1$colorF    278.404    19.651   14.168 < 2e-16 ***
train.data1$colorG     97.087    18.965    5.119 3.08e-07 ***
train.data1$colorH   -349.722    20.597  -16.979 < 2e-16 ***
train.data1$colorJ  -1615.792    30.377  -53.191 < 2e-16 ***
train.data1$clarityIF  1929.624    37.014   52.132 < 2e-16 ***
train.data1$claritySI2  -710.606    19.555  -36.339 < 2e-16 ***
train.data1$clarityVS1  1062.737    20.307   52.335 < 2e-16 ***
train.data1$clarityVS2    815.480    17.968   45.385 < 2e-16 ***
train.data1$clarityVVS1  1557.832    27.573   56.498 < 2e-16 ***
train.data1$clarityVVS2  1529.518    24.112   63.433 < 2e-16 ***
train.data1$depth     -85.513     4.678  -18.279 < 2e-16 ***
train.data1$table     -69.027     3.445  -20.035 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1284 on 43129 degrees of freedom
Multiple R-squared:  0.8964,    Adjusted R-squared:  0.8964
F-statistic: 2.196e+04 on 17 and 43129 DF,  p-value: < 2.2e-16

```

Classification

- Preprocessing:

	carat	cut	color	clarity	depth	table	price
0.16632017	Premium	F	vvs1	0.5000000	0.3269231	Expensive	
0.03534304	Ideal	E	vs2	0.5444444	0.2500000	Inexpensive	
0.06444906	Good	G	si2	0.5611111	0.2692308	Inexpensive	
0.14553015	Premium	G	vs2	0.5500000	0.2884615	Expensive	
0.06237006	Premium	F	vs2	0.5194444	0.3269231	Inexpensive	
0.02286902	Premium	F	vs2	0.5277778	0.3269231	Inexpensive	

- Knn

Preprocessing:

```

> head(data)
  carat cutGood cutIdeal cutPremium cutVery Good colorE colorF colorG colorH colorI colorJ clarityIF claritySI1
1 0.006237006      0      1      0      0      0      1      0      0      0      0      0      0      0
2 0.002079002      0      0      1      0      0      1      0      0      0      0      0      0      1
3 0.006237006      1      0      0      0      0      1      0      0      0      0      0      0      0
4 0.018711019      0      0      1      0      0      0      0      0      0      1      0      0      0
5 0.022869023      1      0      0      0      0      0      0      0      0      0      1      0      0
6 0.008316008      0      0      0      0      1      0      0      0      0      0      0      1      0
  claritySI2 clarityVS1 clarityVS2 clarityVVS1 clarityVVS2 depth table price x y
1      1      0      0      0      0      0.5138889 0.2307692 Inexpensive 0.3677840 0.06757216
2      0      0      0      0      0      0.4666667 0.3461538 Inexpensive 0.3621974 0.06519525
3      0      1      0      0      0      0.3861111 0.4230769 Inexpensive 0.3770950 0.06910017
4      0      0      1      0      0      0.5388889 0.2884615 Inexpensive 0.3910615 0.07181664
5      1      0      0      0      0      0.5638889 0.2884615 Inexpensive 0.4040968 0.07385399
6      0      0      0      0      1      0.5500000 0.2692308 Inexpensive 0.3668529 0.06723260
  z
1 0.07641509
2 0.07264151
3 0.07264151
4 0.08270440
5 0.08647799
6 0.07798742

```

Splitting of data:

```

#hold out evaluation
knndata <- knndata[sample(nrow(knndata)),]
select.dataknn <- sample(1:nrow(knndata), 0.8*nrow(knndata))
train.dataknn <- knndata[select.dataknn,]
test.dataknn <- knndata[-select.dataknn,]
test.knn <- test.dataknn
train.knn <- train.dataknn
head(train.knn)

nrow(test.svm)
nrow(train.svm)

train.knn$price<-NULL
test.knn$price<-NULL
train.def <- train.dataknn$price
test.def <- test.dataknn$price

library(class)
knn.133 <- knn(train.knn, test.knn, train.def, k=133)
knn.155 <- knn(train.knn, test.knn, train.def, k=155)
knn.211 <- knn(train.knn, test.knn, train.def, k=211)
#install.packages("Metrics", dependencies = TRUE)
library(Metrics)
accuracy(test.def, knn.133)
accuracy(test.def, knn.155)
accuracy(test.def, knn.211)

```

Accuracy:

```

> train.def <- train.dataknn$price
> test.def <- test.dataknn$price
> knn.133 <- knn(train.knn, test.knn, train.def, k=133)
Error in knn(train.knn, test.knn, train.def, k = 133) :
  could not find function "knn"
> library(class)
Warning message:
package 'class' was built under R version 3.6.3
> knn.133 <- knn(train.knn, test.knn, train.def, k=133)
> knn.155 <- knn(train.knn, test.knn, train.def, k=155)
> knn.211 <- knn(train.knn, test.knn, train.def, k=211)
> #install.packages("Metrics", dependencies = TRUE)
> library(Metrics)

Attaching package: 'Metrics'

The following objects are masked from 'package:caret':

  precision, recall

Warning message:
package 'Metrics' was built under R version 3.6.3
> accuracy(test.def, knn.133)
[1] 0.9013719
> accuracy(test.def, knn.155)
[1] 0.890063
> accuracy(test.def, knn.211)
[1] 0.8627178
>

```

- Naïve bayes:

Preprocessing:

	carat	cut	color	clarity	depth	table	price
0.16632017	Premium	F	VVS1	0.5000000	0.3269231	Expensive	
0.03534304	Ideal	E	VS2	0.5444444	0.2500000	Inexpensive	
0.06444906	Good	G	SI2	0.5611111	0.2692308	Inexpensive	
0.14553015	Premium	G	VS2	0.5500000	0.2884615	Expensive	
0.06237006	Premium	F	VS2	0.5194444	0.3269231	Inexpensive	
0.02286902	Premium	F	VS2	0.5277778	0.3269231	Inexpensive	

V

output:

```
> library(e1071)
> naive <- naiveBayes(train.naive$price~. , data = train.naive)
> naive

Naive Bayes Classifier for Discrete Predictors

Call:
naiveBayes.default(x = X, y = Y, laplace = laplace)

A-priori probabilities:
Y
Inexpensive    Expensive
  0.6404338    0.3595662

Conditional probabilities:
      carat
Y      [,1]      [,2]
Inexpensive 0.06500432 0.04351133
Expensive   0.22995684 0.07872061

      cut
Y      Fair      Good      Ideal      Premium      Very Good
Inexpensive 0.02789839 0.08865248 0.44510783 0.21989434 0.21844695
Expensive   0.03280485 0.09706110 0.31908997 0.31419180 0.23685228

      color
Y      D      E      F      G      H      I      J
Inexpensive 0.14303807 0.21218700 0.18765378 0.20791721 0.12954118 0.08228398 0.03737878
Expensive   0.09383862 0.12728796 0.15957721 0.21203919 0.19670018 0.13334622 0.07721062

      clarity
Y      I1      IF      SI1      SI2      VS1      VS2      VVS1      VVS2
Inexpensive 0.01429295 0.04396439 0.22155884 0.12986684 0.15888696 0.23035172 0.08999132 0.11108699
Expensive   0.01282547 0.01688580 0.27597319 0.24419954 0.13527971 0.22409126 0.02777778 0.06296726

      depth
Y      [,1]      [,2]
Inexpensive 0.5207540 0.03861706
Expensive   0.5210487 0.04150635

      table
Y      [,1]      [,2]
Inexpensive 0.2735346 0.04291354
Expensive   0.2856311 0.04202906
```

Accuracy:

```
> confusionMatrix(table(pre_naive,test.naive$price))
Confusion Matrix and Statistics

pre_naive      Inexpensive Expensive
Inexpensive      6324         211
Expensive         601        3652

      Accuracy : 0.9247
      95% CI   : (0.9196, 0.9296)
      No Information Rate : 0.6419
      P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.8398

      Mcnemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.9132
      Specificity : 0.9454
      Pos Pred Value : 0.9677
      Neg Pred Value : 0.8587
      Prevalence : 0.6419
      Detection Rate : 0.5862
      Detection Prevalence : 0.6058
      Balanced Accuracy : 0.9293

      'Positive' Class : Inexpensive
```


- Support vector machines:

Preprocessing:

```

carat cutGood cutIdeal cutPremium cutVery Good colorE colorF colorG colorH colorI colorJ clarityIF
0.10602911 0 0 0 0 1 0 0 1 0 0 0 0
0.22245322 0 0 0 1 0 0 0 0 1 0 0 0
0.04158004 0 0 1 0 0 0 1 0 0 0 0 0
0.18918919 0 0 1 0 0 0 0 0 1 0 0 0
0.10602911 0 0 0 0 1 0 0 0 0 0 1 0
0.04573805 0 0 1 0 0 1 0 0 0 0 0 0
claritySI1 claritySI2 clarityVS1 clarityVS2 clarityVVS1 clarityVVS2 depth table price
0 1 0 0 0 0 0.4444444 0.2884615 Inexpensive
0 0 0 0 1 0 0.5027778 0.2884615 Expensive
0 0 1 0 0 0 0.4861111 0.3269231 Inexpensive
1 0 0 0 0 0 0.4388889 0.3076923 Expensive
0 1 0 0 0 0 0.4972222 0.2307692 Inexpensive
1 0 0 0 0 0 0.5222222 0.2884615 Inexpensive

```

Model:

```

call:
svm(formula = train.svm$price ~ ., data = train.svm, kernel = "linear",
    cost = 0.1, scale = F)

Parameters:
  SVM-Type:  c-classification
  SVM-Kernel: linear
    cost: 0.1

Number of Support Vectors: 12203

( 6102 6101 )

Number of Classes: 2

Levels:
Inexpensive Expensive

```

Accuracy:

```

> confusionMatrix(table(pre_svm,test.svm$price))
Confusion Matrix and Statistics

pre_svm      Inexpensive Expensive
Inexpensive      6661      118
Expensive        238      3771

      Accuracy : 0.967
      95% CI   : (0.9635, 0.9703)
 No Information Rate : 0.6395
 P-Value [Acc > NIR] : < 2.2e-16

      Kappa   : 0.9289

McNemar's Test P-Value : 2.845e-10

      Sensitivity : 0.9655
      Specificity : 0.9697
      Pos Pred Value : 0.9826
      Neg Pred Value : 0.9406
      Prevalence : 0.6395
      Detection Rate : 0.6174
      Detection Prevalence : 0.6284
      Balanced Accuracy : 0.9676

      'Positive' Class : Inexpensive

```

- Random forest:

Preprocessing:

```

> rfdata$price <- cost
> head(rfdata)
  X carat      cut color clarity depth table price    x    y    z
1 1  0.23    Ideal     E    SI2   61.5    55     0 3.95 3.98 2.43
2 2  0.21  Premium     E    SI1   59.8    61     0 3.89 3.84 2.31
3 3  0.23     Good     E    VS1   56.9    65     0 4.05 4.07 2.31
4 4  0.29  Premium     I    VS2   62.4    58     0 4.20 4.23 2.63
5 5  0.31     Good     J    SI2   63.3    58     0 4.34 4.35 2.75
6 6  0.24 Very Good     J   VVS2   62.8    57     0 3.94 3.96 2.48
>

```

Model:

```

call:
  randomForest(formula = train.rf$price ~ ., data = train.rf)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 2

      OOB estimate of  error rate: 2.35%
Confusion matrix:
      0      1 class.error
0 26984   582  0.02111297
1   431 15155  0.02765302
>

```

Accuracy:

```

Confusion Matrix and Statistics

pre_rf      0      1
      0 6854  118
      1  141 3675

      Accuracy : 0.976
      95% CI   : (0.9729, 0.9788)
No Information Rate : 0.6484
P-Value [Acc > NIR] : <2e-16

      Kappa   : 0.9474

McNemar's Test P-Value : 0.1716

      Sensitivity : 0.9798
      Specificity : 0.9689
      Pos Pred Value : 0.9831
      Neg Pred Value : 0.9631
      Prevalence : 0.6484
      Detection Rate : 0.6353
      Detection Prevalence : 0.6463
      Balanced Accuracy : 0.9744

      'Positive' Class : 0

```

Each forest evaluation:

```

> gettree(rf, 1, labelvar = τ)
  left daughter right daughter split var split point status prediction
1          2          3      carat  0.15904366      1      <NA>
2          4          5      table  0.24903846      1      <NA>
3          6          7      table  0.34326923      1      <NA>
4          8          9      carat  0.13409563      1      <NA>
5         10         11 clarity  52.00000000      1      <NA>
6         12         13      cut   1.00000000      1      <NA>
7         14         15 clarity   9.00000000      1      <NA>
8         16         17 clarity  63.00000000      1      <NA>
9         18         19      cut  11.00000000      1      <NA>
10        20        21      cut  13.00000000      1      <NA>
11        22        23      depth  0.44583333      1      <NA>
12        24        25      color  34.00000000      1      <NA>
13        26        27 clarity   1.00000000      1      <NA>
14        28        29 clarity   1.00000000      1      <NA>
15        30        31      cut   1.00000000      1      <NA>
16        32        33      table  0.22980769      1      <NA>
17        34        35      carat  0.11538462      1      <NA>
18        36        37      carat  0.15072765      1      <NA>
19        38        39      color  31.00000000      1      <NA>
20        40        41      color  31.00000000      1      <NA>
21        42        43      cut   2.00000000      1      <NA>
22        44        45      color  2.00000000      1      <NA>
23        46        47      carat  0.14033264      1      <NA>
24        48        49      depth  0.40972222      1      <NA>
25        50        51      carat  0.20893971      1      <NA>
26        52        53      carat  0.22557173      1      <NA>
27        54        55      carat  0.17983368      1      <NA>
28        56        57      carat  0.24324324      1      <NA>

```